SherdMatch

AI-Based Archaeological Pottery Fragment Matching

From a plan to a working, evaluated system: appearance + fracture-edge geometry on the Daxinzhuang-18K benchmark

PRT LLT · 21CSE451T

Review 2

Vansh Agrawal · RA2311026010120
CINTEL · Section AC-1
Faculty: Dr Nithyashri J



PROGRESS SINCE REVIEW 1

From plan to evidence

REVIEW 1

Dataset named, schema to verify, pipeline planned, no results

REVIEW 2

Data acquired and verified, pipeline built, 5 experiments, working demo

124
real sherds in the test set

111
expert-confirmed joins

16
model configurations compared

5
experiments run (E1–E5)

- **Data**: Daxinzhuang-18K found on Zenodo; only the needed parts fetched (range requests, CRC-checked)
- **Pipeline**: segmentation, CLAHE, contours, colour/LBP/ResNet-50/DINOv2 features
- **Geometry**: complementary-curvature edge matcher with Kabsch + ICP alignment
- **Model**: logistic-regression fusion, compared with RF and SVM
- **Evaluation**: leak-free vessel-grouped CV, bootstrap CIs, ablations
- **Prototype**: Gradio demo with ranked candidates and fit preview

DATASET

Daxinzhuang-18K, verified

Real photos	36,532 photos · 17,635 records	metadata + photo segmentation
Gold standard	124 sherds · 111 expert joins + poses	main train / test set
Simulated	50,736 synthetic break pairs	geometry sanity check (2,212 used)
Refit sheet	251 expert joining groups	ground truth, raw-photo test (next)

Zenodo 21553799 · Ji et al., 2026 · CC-BY 4.0

1.46%

of the 7,626 sherd pairs are joins: extreme class imbalance

33

vessel groups (2–11 sherds each), used to split train/test





METADATA UNDERSTANDING

What pixels cannot tell us

PHOTO ID ≠ SHERD ID

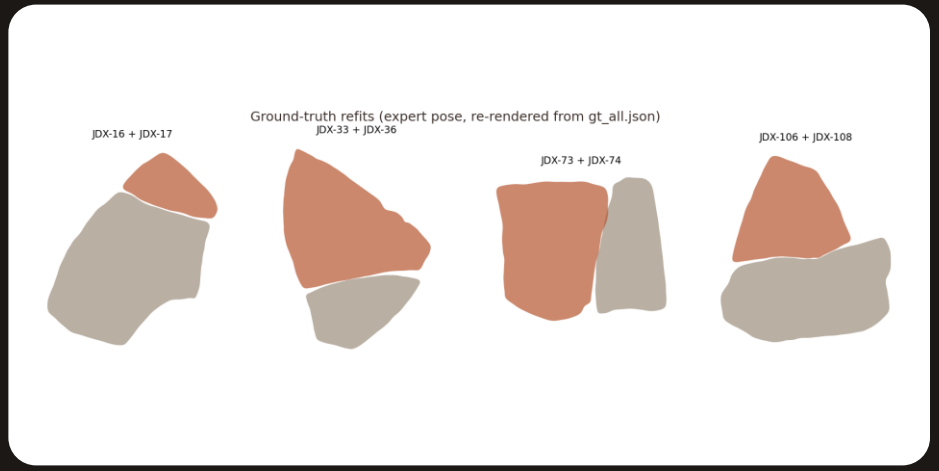
Sherd JD00001 → photos JD_00001 (exterior) + JD_00002 (interior). Mapping by number would pair the wrong photos. Fixed via the Exterior_Photo column.

POSE CONVENTIONS VERIFIED

Gold: recentre on alpha centroid, rotate, place. Simulated: image centre, opposite sign. Verified: true pairs touch 75–170 px with ≤7% overlap.

USED FOR

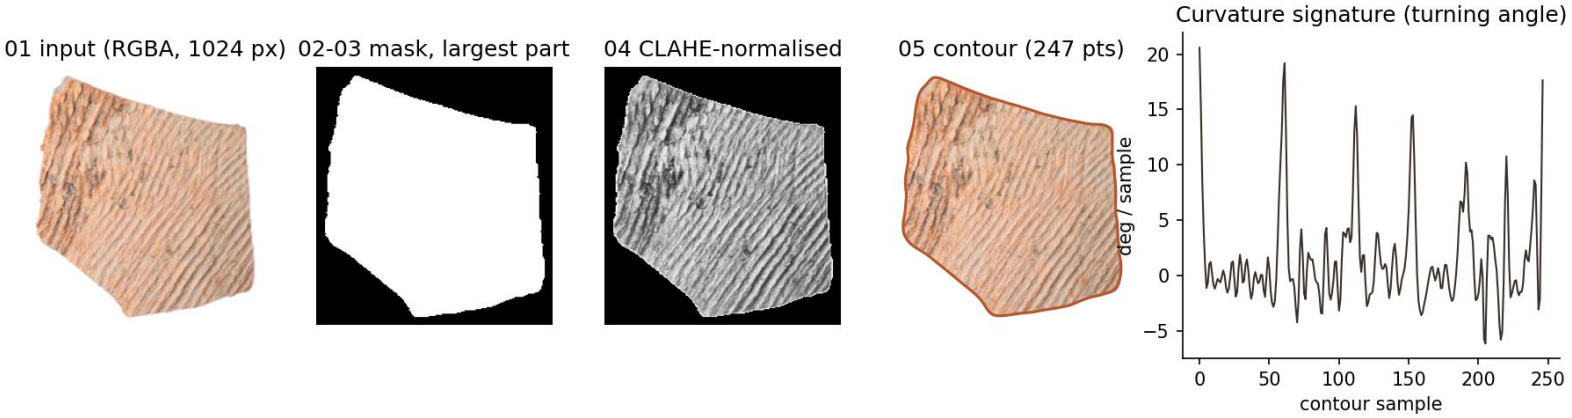
- **Vessel groups** → leak-free folds
- **Notes** (spalling) explain errors
- **Layer, part** → future priors
- Only 547 of 17,635 records typed: too sparse as a label

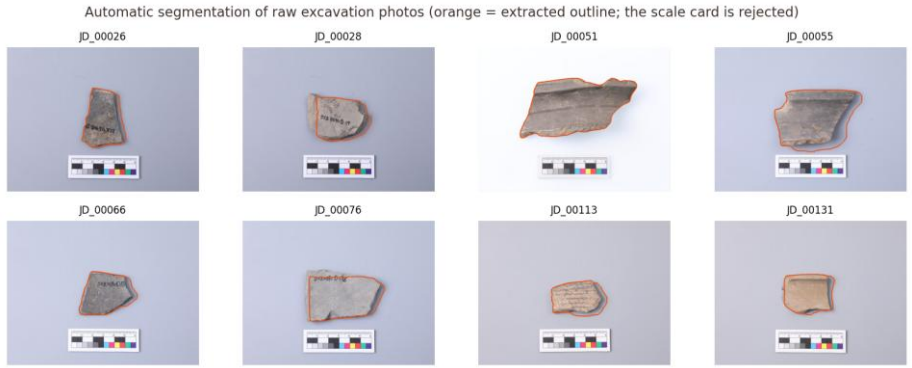


Expert joins re-rendered with our own pose code. Visible gaps along the joints are edge erosion, the central obstacle for geometry.

THE PROCESS, IMPLEMENTED

Preprocessing pipeline

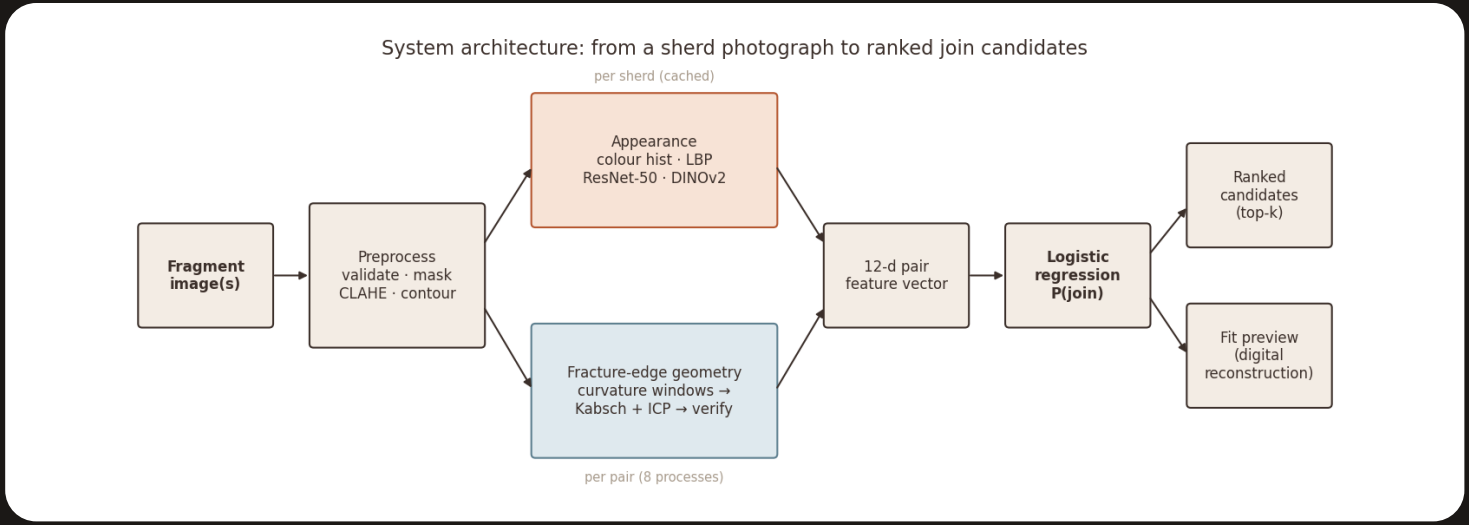




- Scale kept fixed: edge lengths must match
- Raw photos: Otsu → card filter → **GrabCut**
- Remaining error: light edges cut, shadows kept

MODEL SELECTION

Two cues, one interpretable classifier



4 appearance: colour histogram, LBP texture, ResNet-50, DINOv2

8 geometry: contact, contact fraction, overlap, rmse, fit score, stripe-angle gap (+confidence), area ratio

PATTERN TO RECOGNISE

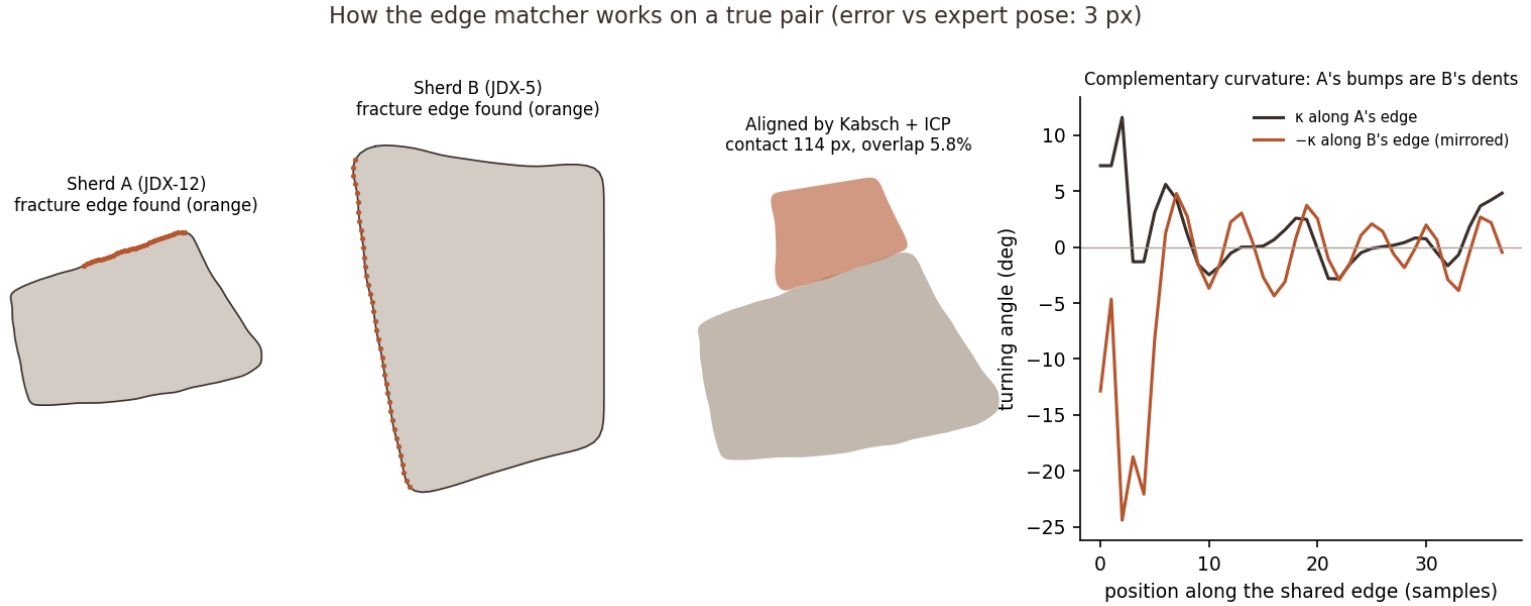
Does sherd pair (A, B) join? Used as ranking: sort the gallery by P(join).

LOGISTIC REGRESSION

- $P(\text{join}) = \sigma(w^T z + b)$
- 12 standardised pair features
- class-balanced, L2 (C = 1)

FRACTURE-EDGE GEOMETRY

Complementary curvature matching



01 Signature

turning angle every 3 px

02 Seeds

A vs reversed −B windows

03 Align

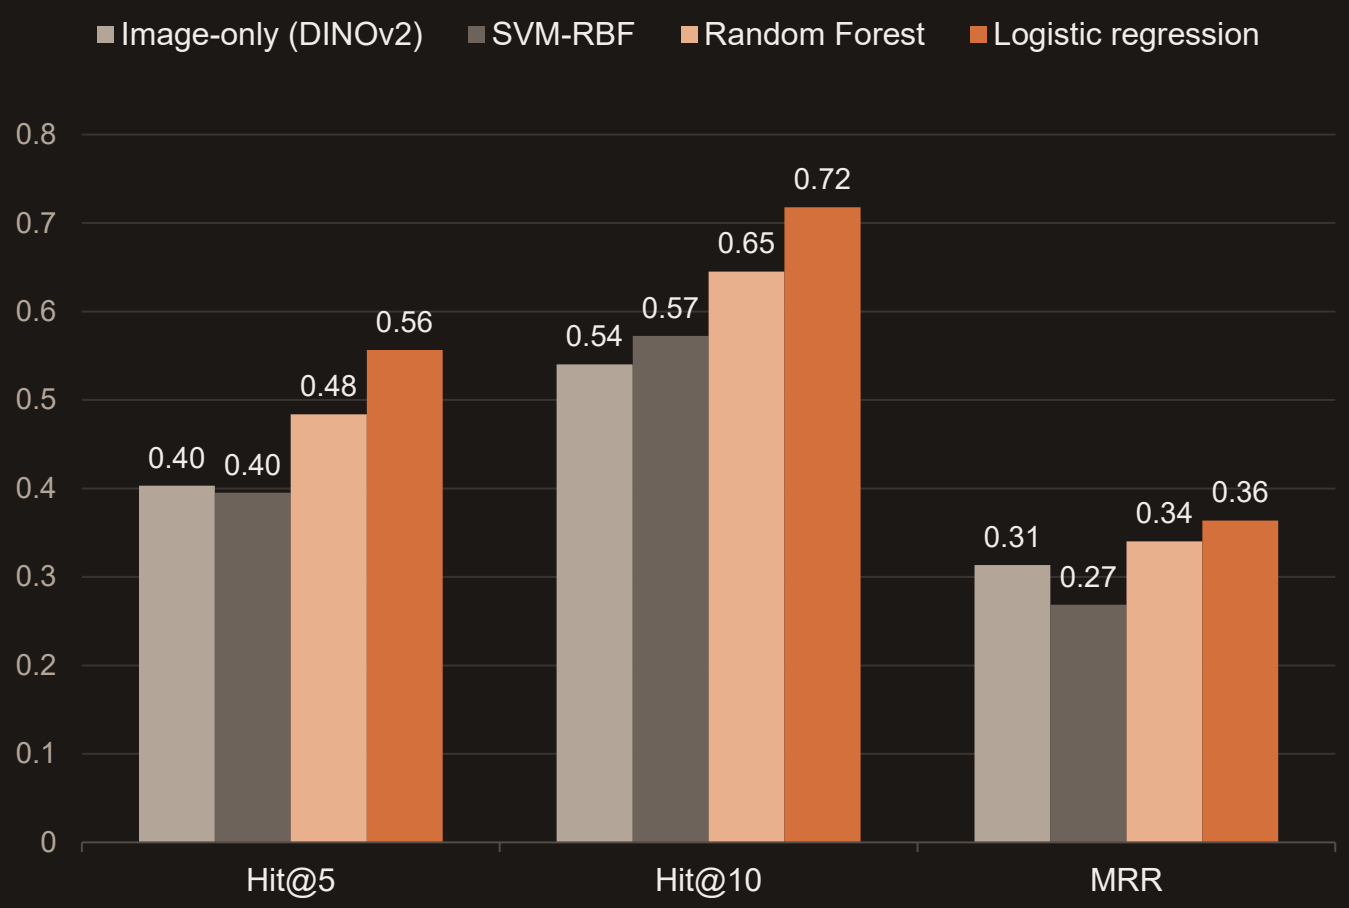
Kabsch, then ICP

04 Verify

contact ↑ overlap ↓ stripes agree

JUSTIFICATION

Why logistic regression



DATA SIZE

111 positives fit a 13-parameter model; non-linear models and deep nets overfit

INTERPRETABILITY

Signed weight per cue: experts can see why a sherd was suggested

COST AND RESULT

Milliseconds to train; best Hit@5 and Hit@10 of all classifiers

EXPERIMENTAL SETUP

Leak-free evaluation

VALIDATION

- 5-fold **GroupKFold** over the 33 vessel groups
- No vessel in both train and test
- Each query ranked against all 123 others
- F1 threshold from inner CV on training pairs

METRICS AND WHY

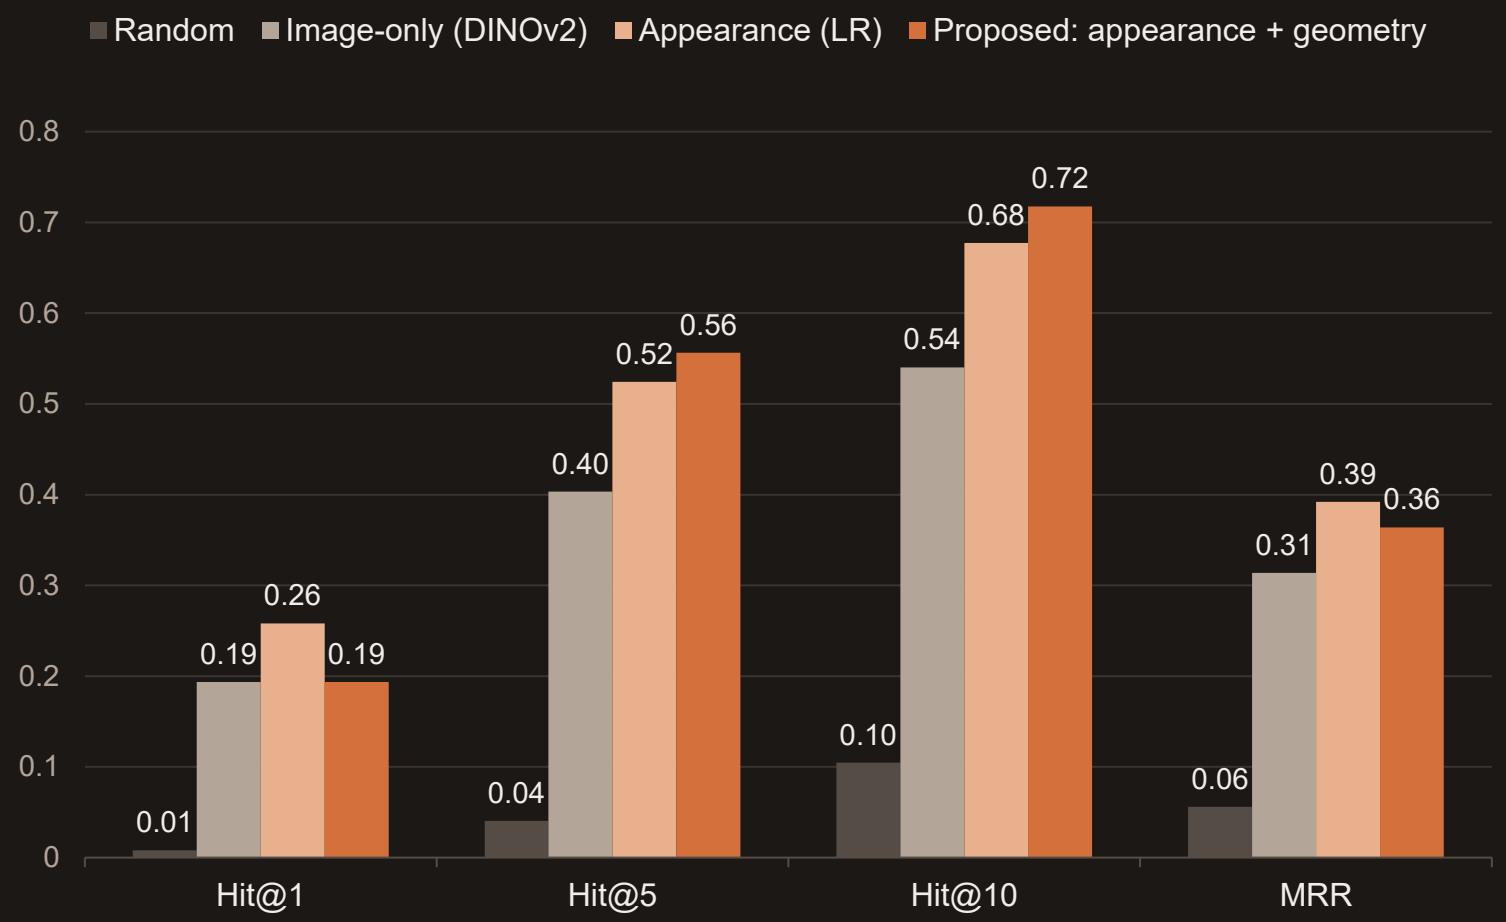
- **Hit@k**: what the expert checks
- **MRR, mAP**: rank quality
- **P / R / F1, PR-AUC**: pair decisions
- Accuracy is misleading at 1.46% positives
- **Bootstrap 95% CIs** (5,000×)

EXPERIMENTS

- E1 image-only vs fusion
- E2 LR vs RF vs SVM
- E3 ablation of every cue
- E4 geometry on synthetic breaks
- E5 pose accuracy on real joins
- E6 raw photos (next step)

RESULTS (E1)

Better than image-only retrieval



72%

of sherds have a true partner in the top 10

+17.7 pts

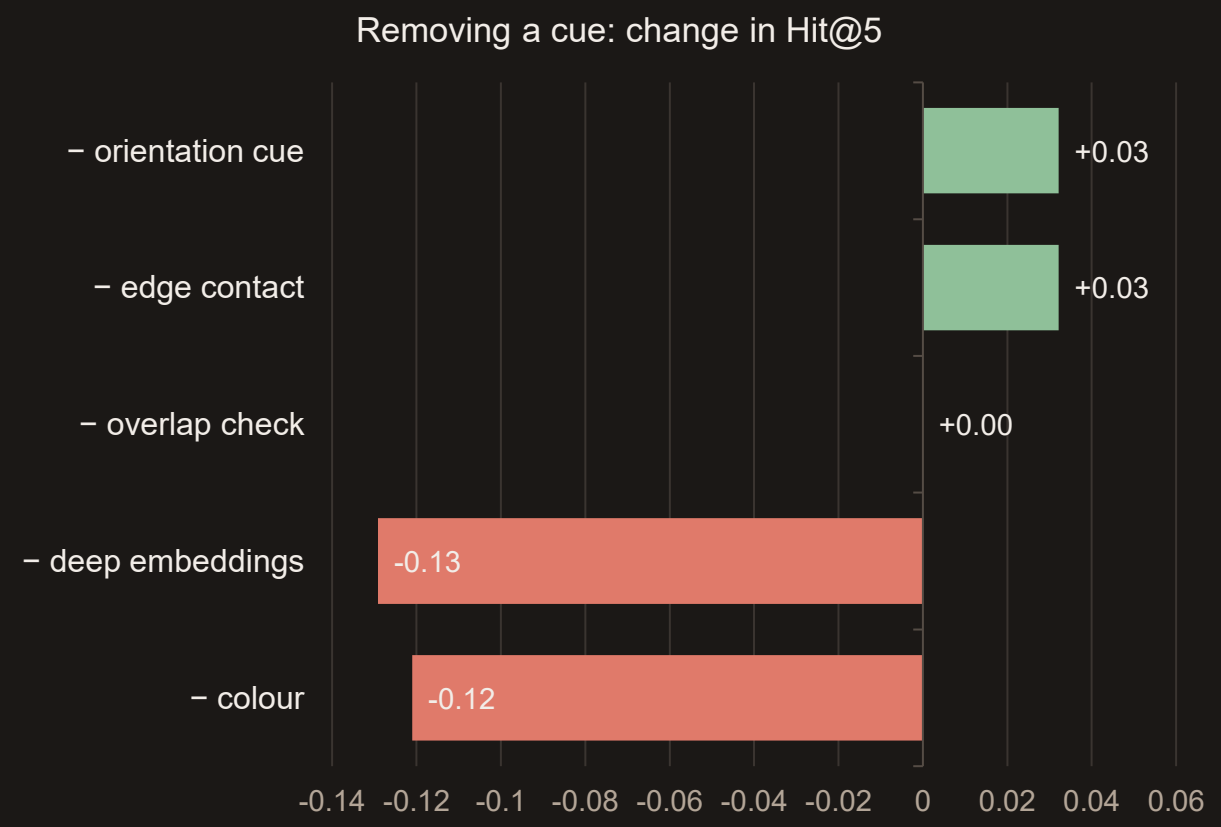
Hit@10 vs image-only (95% CI +9.7 to +26.6)

+15.3 pts

Hit@5 vs image-only (95% CI +5.6 to +25.0)

TESTING THE HYPOTHESIS

Does geometry pull its weight?



Red = the cue helped (removing it hurts). Green = removing it helps.

Hit@1	**-6.5**	95% CI -12.1 to -1.6 significant
Hit@5	**+3.2**	95% CI -2.4 to +9.7 not significant
Hit@10	**+4.0**	95% CI -2.4 to +10.5 not significant

Adding geometry to learned appearance: **H1 only partly supported**

DIAGNOSIS (E4, E5)

The bottleneck is the edge, not the matcher

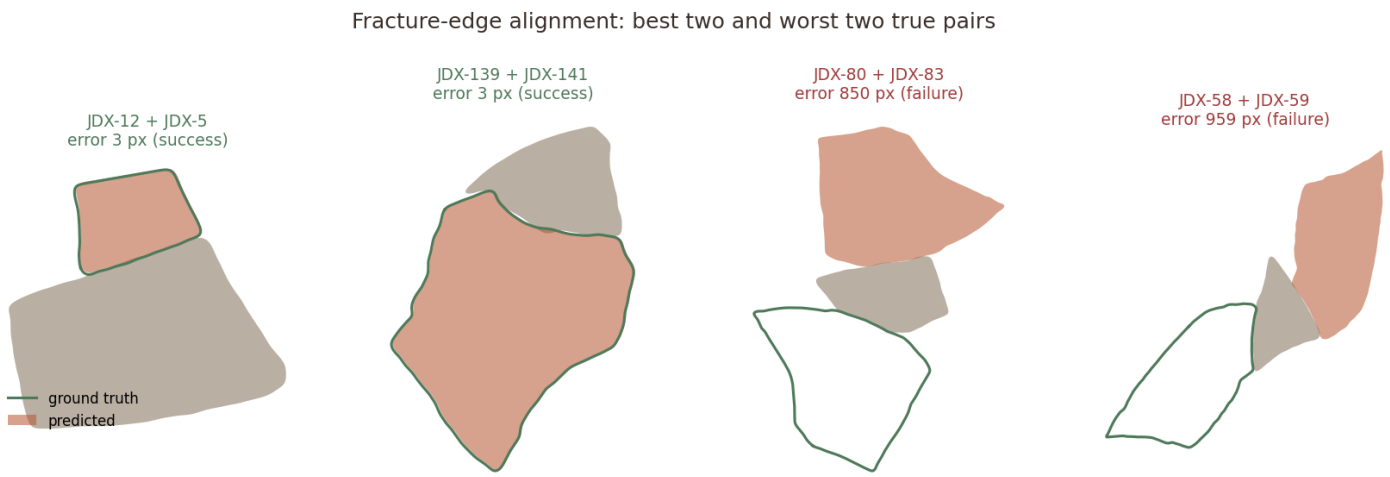
47%

clean synthetic breaks aligned within 50 px

17%

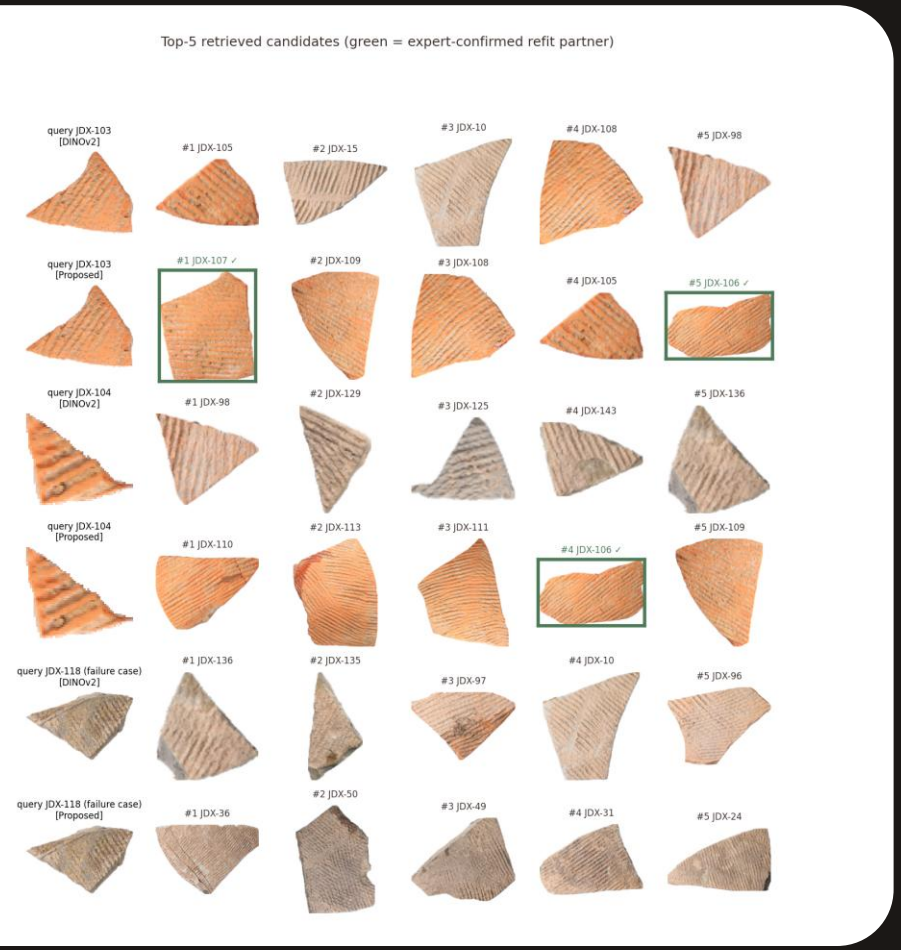
real eroded sherds aligned within 50 px

- At the expert pose, true partners touch along **87 px** (median)
- The best **wrong** pose found by ICP touches along **129 px**
- On worn edges, maximising contact chooses wrong fits



RESULTS ANALYSIS

Where it fails, and why



11%

of wrong top-1 picks are from the same vessel: the confusion is between vessels of one ware

7×

more same-vessel sherds in the top 5 than chance

0.44→0.71

Hit@5, smallest vs largest third of sherds

3.4×

PR-AUC over chance (precision 0.37, recall 0.23, F1 0.29)

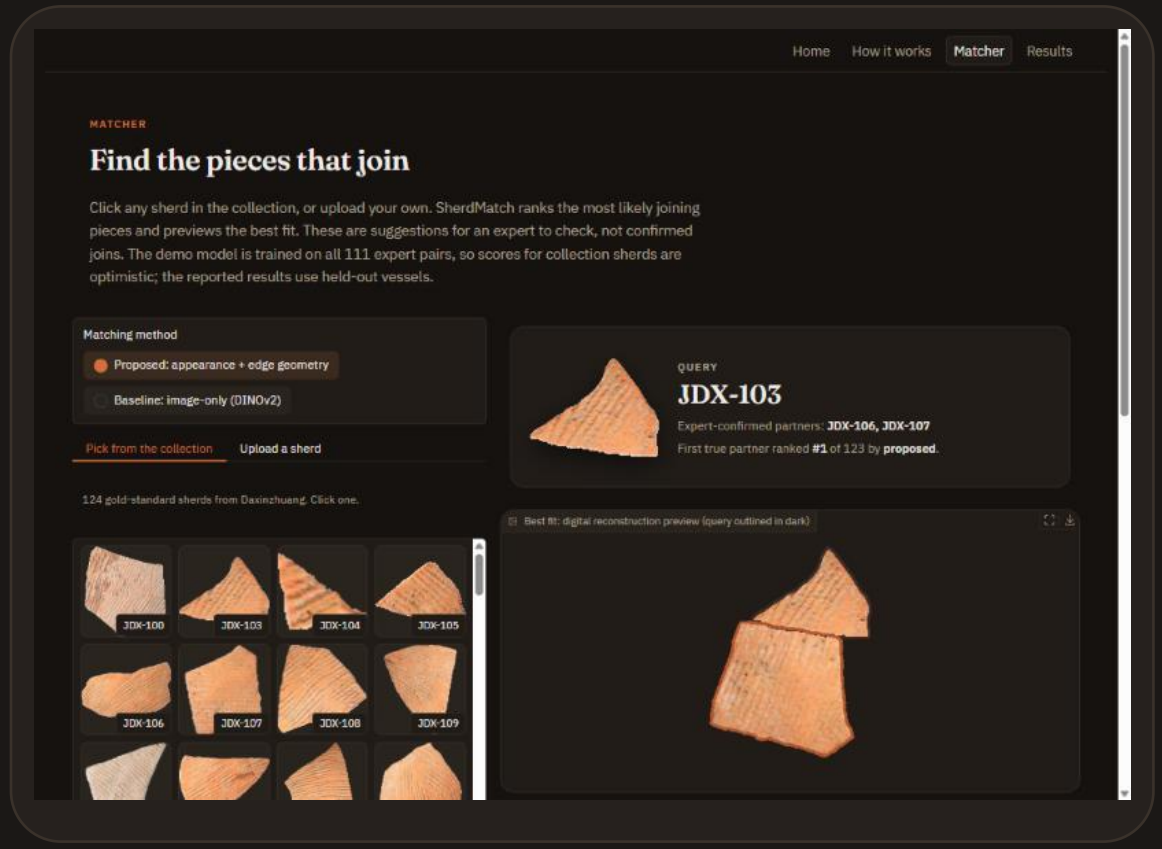
Accuracy is 91.6%, but predicting 'no join' for everything already scores 92.8% on these test pairs, so it is not a useful measure here.

PROTOTYPE

SherdMatch web app



HOME: the problem, the approach, the headline results



MATCHER: click a sherd, get ranked joins and a fit preview

4 pages: Home · How it works · Matcher · Results | JDX-103: true partner ranked #1, vs #28 with image-only | python app.py · github.com/vansh070605/SherdMatch

INFERENCE

What the evidence supports

USEFUL TRIAGE TOOL

An expert checking 10 suggestions finds a partner for 72% of sherds

H1 PARTLY SUPPORTED

Better than image-only, but the gain comes from fusing appearance cues

CLAIMS WE DO NOT MAKE

- Not an automatic reconstruction
- Not 'geometry is useless': 3-D and thickness untested
- One site, one ware

EDGES ARE THE BOTTLENECK

2-D outline geometry fails on eroded edges, not on clean ones

COLOUR MATTERS MOST

Vessels fired unevenly: colour identifies the vessel better than shape

LIMITATIONS & FUTURE SCOPE

What comes next

LIMITATIONS

- 2-D outlines miss thickness and wall curvature
- Max-contact pose choice fails on worn edges
- Only 111 positives: wide confidence intervals
- Exterior view only; one site and ware
- Raw-photo masks: shadows, light edges

Interior view + edge-band colour — LOW

independent evidence across the joint

Learned pose scoring — MEDIUM

train on expert vs false ICP poses

Contrastive fine-tuning — MEDIUM

better vessel grouping

Multi-piece assembly — HIGH

erosion-aware simulation → full reconstruction

REFERENCES

Sources and code

Ji, D., et al. (2026). Daxinzhuang-18K: A large-scale 2D benchmark for pottery sherd retrieval and reassembly. Zenodo. doi:10.5281/zenodo.21553799

Freeman, H., & Garder, L. (1964). Apictorial jigsaw puzzles. IEEE Trans. Electronic Computers, EC-13(2), 118–127.

Huang, Q.-X., et al. (2006). Reassembling fractured objects by geometric matching. ACM TOG, 25(3), 569–578.

Pinho, M., Mota, G., & da Costa, G. (2025). A deep learning approach to assist in pottery reconstruction from its sherds. Heritage, 8(5), 167.

Yu, X., et al. (2026). Mitigating spurious features by contrastive learning in pottery sherd recognition. npj Heritage Science, 14, 135.

Oquab, M., et al. (2024). DINOv2: Learning robust visual features without supervision. TMLR.

He, K., et al. (2016). Deep residual learning for image recognition. CVPR.

Kabsch, W. (1976). Acta Cryst. A32, 922–923 · Besl, P., & McKay, N. (1992). IEEE TPAMI 14(2), 239–256.

Full list (23 references): Review 2 case study report, Section 16.

Code, results, report

github.com/vansh070605/SherdMatch

Thank you

Questions?

